



CSE422: Artificial Intelligence

Project Title: Road Accident Severity Prediction

Lab Section: 10

Group: 10

Spring 2024

ID	Name
21301470	Syeda Ifroza Ahmed
21301196	Swachcha Podder

Table of Contents

Section No.	Content	Page No.
1	Introduction	2
2	Dataset Description	3
3	Dataset Pre-processing	6
4	Feature Scaling	10
5	Dataset Splitting	10
6	Model Training & Testing	10
7	Model Selection/ Comparison Analysis	12
8	Conclusion	28

INTRODUCTION

Road accident severity prediction

The road accident severity prediction aims to develop a robust predictive model that can accurately assess the severity of road accidents. By analyzing various features such as Driving_experience, Area_accident_occured, Weather_conditions, and Number_of_casualties, etc, the project seeks to anticipate the potential severity of accidents.

The problem it is aiming to solve

The project addresses the significant issue of road accidents and their associated consequences, including fatalities. By developing a predictive model for road accident severity, the project aims to provide stakeholders with valuable insights to proactively mitigate the impact of accidents and improve road safety measures. The ultimate goal is to provide valuable insights to stakeholders such as law enforcement agencies, transportation authorities, and emergency responders, enabling them to allocate resources more effectively, implement preventive measures, and improve overall road safety.

Motivation Behind the Project

The motivation stems from the urgent need to enhance road safety and reduce the human and economic toll of road accidents. With road accidents posing a continuous threat to public safety worldwide, there's a pressing need for proactive measures to address this issue. By leveraging machine learning models and techniques to predict accident severity, the project aims to empower stakeholders with the tools and information needed to prevent accidents, optimize emergency response efforts, and ultimately save lives.

DATASET DESCRIPTION

Source: Kaggle

Link: <https://www.kaggle.com/datasets/kanuriviveknag/road-accidents-severity-dataset>

Reference: Vivek Nag Kanuri

Dataset description

1. How many features:

There are 10 features initially. Such as: Day_of_week, Age_band_of_driver, Sex_of_driver, Educational_level, Vehicle_driver_relation, Driving_experience, Type_of_vehicle, Owner_of_vehicle, Service_year_of_vehicle, Defect_of_vehicle, Area_accident_occured, Lanes_or_Medians, Road_allignment, Types_of_Junction, Road_surface_type, Road_surface_conditions, Light_conditions, Weather_conditions, Type_of_collision, Number_of_vehicles_involved, Number_of_casualties, Vehicle_movement, Casualty_class, Sex_of_casualty, Age_band_of_casualty, Casualty_severity, Work_of_casualty, Fitness_of_casualty, Pedestrian_movement, Cause_of_accident, Accident_severity.

The outcome has 3 classes, 'Slight Injury', 'Serious Injury', 'Fatal Injury'

After pre-processing there are 17 features.

2. Is this a classification or regression problem? Why do you think so?

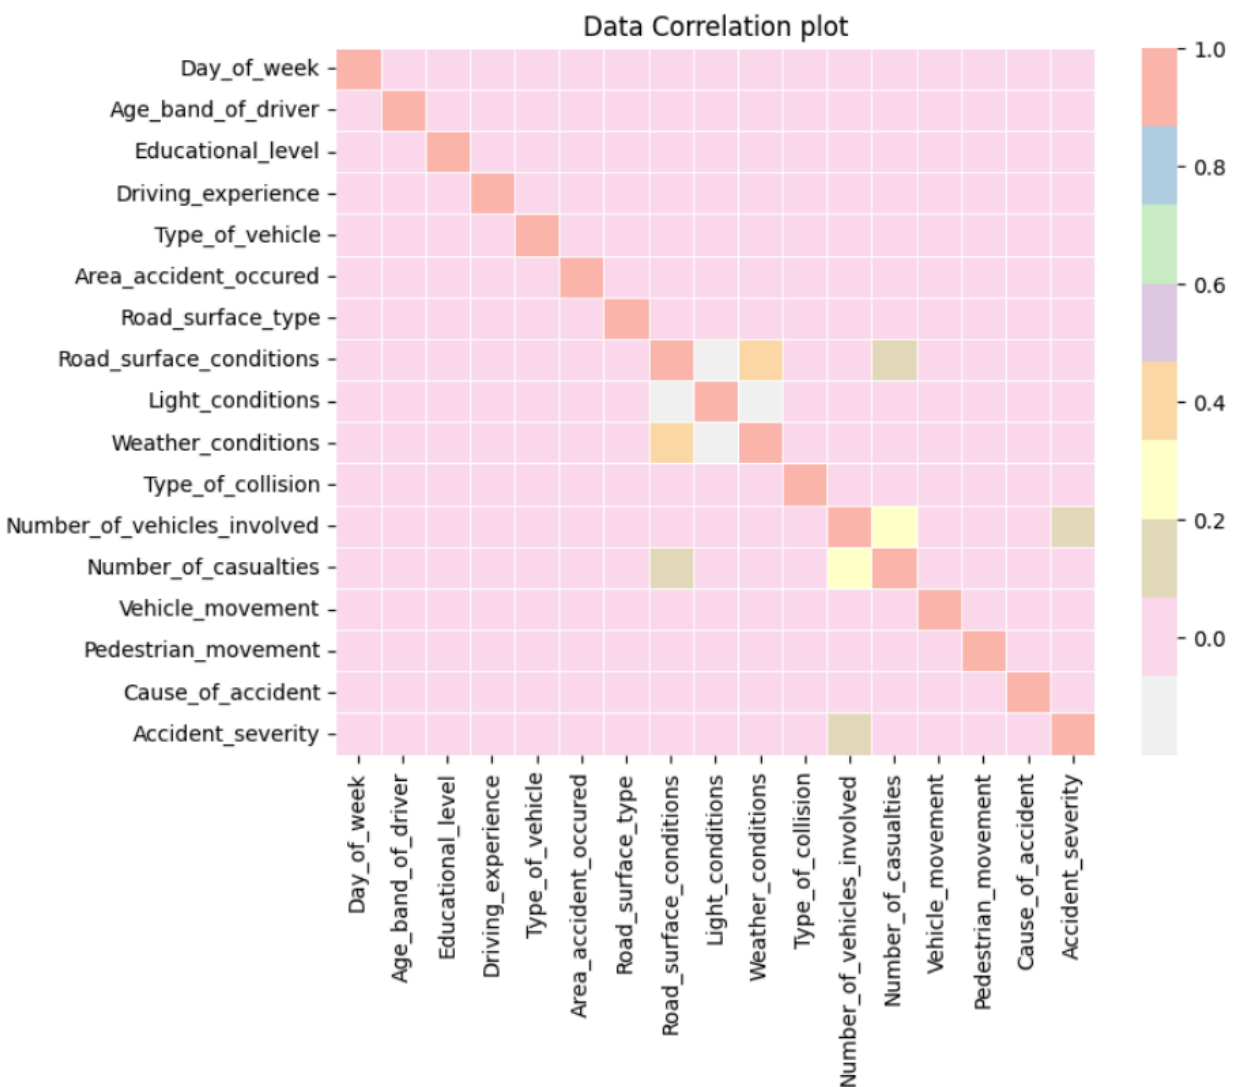
For our project, we are addressing a multivalued classification problem. Classification involves the task of predicting categories or classes based on given input features. In the context of road accident severity prediction, our aim is to predict whether features like Driving experience, Weather conditions affect the accident severity. These outcomes, " Serious Injury " and "Slight Injury", "Fatal Injury" are distinct categories that we are trying to predict based on various accident related features. Therefore, the nature of our task, which involves predicting a categorical label (Serious Injury, Slight Injury, Fatal Injury) signifies that our project falls within the domain of classification problems.

3. How many data points?

Initially 12316

After preprocessing 10932

4. What kind of features are in your dataset? Quantitative or categorical?

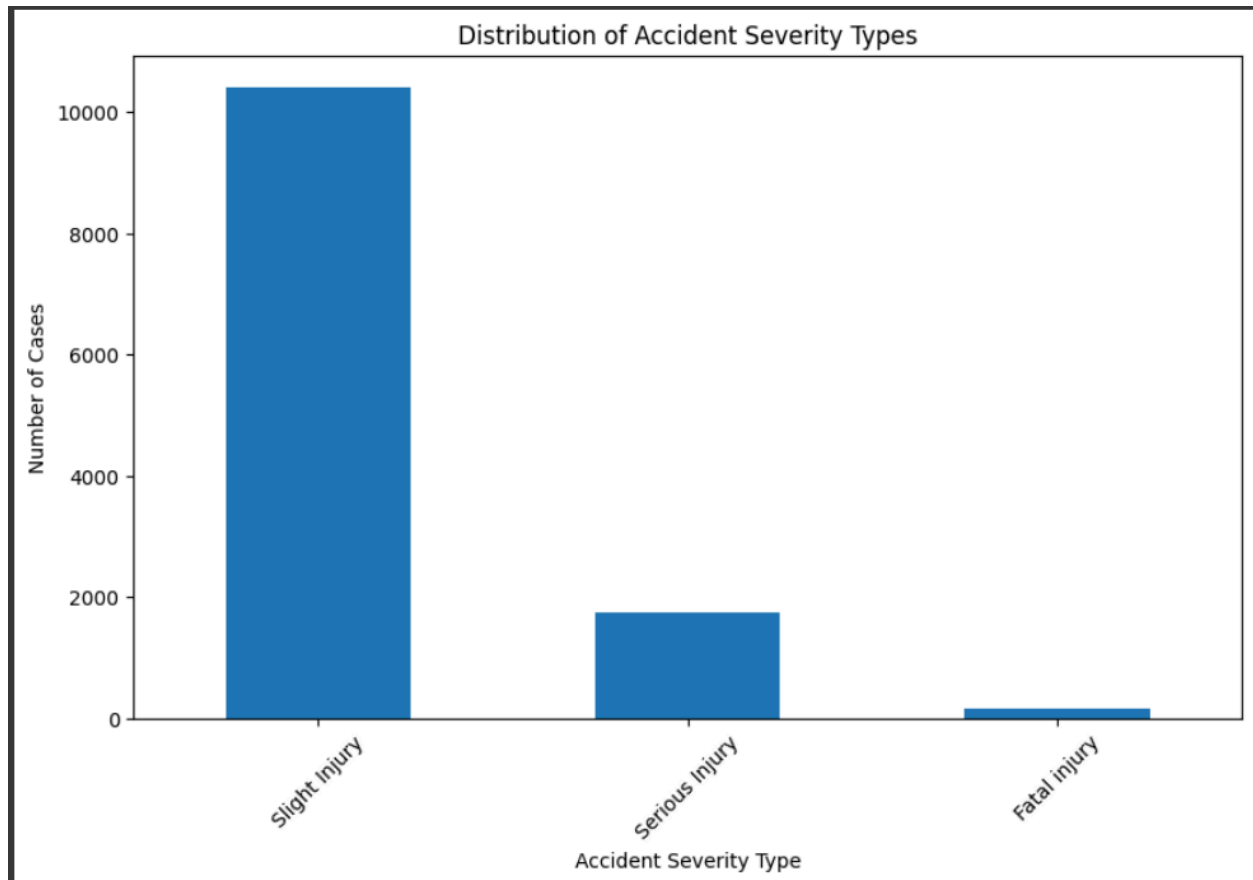


Imbalanced Dataset

1. For the output feature, do all unique classes have an equal number of instances or not?

Ans: No

2. Represent using a bar chart of N classes (N=number of classes you have in your dataset).



DATASET PRE-PROCESSING

Faults: Null values

```
1 road.isnull().sum()
```

```
Time          0
Day_of_week   0
Age_band_of_driver  0
Sex_of_driver  0
Educational_level  741
Vehicle_driver_relation  579
Driving_experience  829
Type_of_vehicle  950
Owner_of_vehicle  482
Service_year_of_vehicle  3928
Defect_of_vehicle  4427
Area_accident_occured  239
Lanes_or_Medians  385
Road_allignment  142
Types_of_Junction  887
Road_surface_type  172
Road_surface_conditions  0
Light_conditions  0
Weather_conditions  0
Type_of_collision  155
Number_of_vehicles_involved  0
Number_of_casualties  0
Vehicle_movement  308
Casualty_class  0
Sex_of_casualty  0
Age_band_of_casualty  0
Casualty_severity  0
Work_of_casualty  3198
Fitness_of_casualty  2635
Pedestrian_movement  0
Cause_of_accident  0
Accident_severity  0
dtype: int64
```

Drop null value columns:

```
road2.drop(['Time', 'Sex_of_driver', 'Defect_of_vehicle',
            'Vehicle_driver_relation', 'Owner_of_vehicle',
```



```
'Service_year_of_vehicle','Lanes_or_Medians','Road_allignment',
',
'Types_of_Junction','Casualty_severity','Casualty_class',
'Sex_of_casualty','Age_band_of_casualty','Work_of_casualty',
'Fitness_of_casualty'],axis=1,inplace=True)
```

Drop null value rows:

```
1 road2 = road2.dropna(axis = 0, subset = ['Vehicle_movement','Type_of_collision','Type_of_vehicle'])
```

Impute values:

```
1 #Imputing values
2
3 from sklearn.impute import SimpleImputer
4
5 imputer = SimpleImputer(missing_values=np.nan, strategy='most_frequent')
6 imputer.fit(road2[['Educational_level']])
7 road2[['Educational_level']] = imputer.transform(road2[['Educational_level']])
8
9 imputer = SimpleImputer(missing_values=np.nan, strategy='most_frequent')
10 imputer.fit(road2[['Driving_experience']])
11 road2[['Driving_experience']] = imputer.transform(road2[['Driving_experience']])
12
13 imputer = SimpleImputer(missing_values=np.nan, strategy='most_frequent')
14 imputer.fit(road2[['Area_accident_occured']])
15 road2[['Area_accident_occured']] = imputer.transform(road2[['Area_accident_occured']])
16
17 imputer = SimpleImputer(missing_values=np.nan, strategy='most_frequent')
18 imputer.fit(road2[['Road_surface_type']])
19 road2[['Road_surface_type']] = imputer.transform(road2[['Road_surface_type']])
20
21 1 road2.isnull().sum()
```

Fault: Categorical values

```

1 road.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12316 entries, 0 to 12315
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Time                                  12316 non-null  object
1   Day_of_week                          12316 non-null  object
2   Age_band_of_driver                   12316 non-null  object
3   Sex_of_driver                        12316 non-null  object
4   Educational_level                    11575 non-null  object
5   Vehicle_driver_relation              11737 non-null  object
6   Driving_experience                   11487 non-null  object
7   Type_of_vehicle                     11366 non-null  object
8   Owner_of_vehicle                    11834 non-null  object
9   Service_year_of_vehicle              8388 non-null   object
10  Defect_of_vehicle                   7889 non-null   object
11  Area_accident_occured               12077 non-null  object
12  Lanes_or_Medians                   11931 non-null  object
13  Road_allignment                    12174 non-null  object
14  Types_of_Junction                 11429 non-null  object
15  Road_surface_type                  12144 non-null  object
16  Road_surface_conditions            12316 non-null  object
17  Light_conditions                   12316 non-null  object
18  Weather_conditions                 12316 non-null  object
19  Type_of_collision                  12161 non-null  object
20  Number_of_vehicles_involved         12316 non-null  int64
21  Number_of_casualties                12316 non-null  int64
22  Vehicle_movement                   12008 non-null  object
23  Casualty_class                     12316 non-null  object
24  Sex_of_casualty                    12316 non-null  object
25  Age_band_of_casualty                12316 non-null  object
26  Casualty_severity                   12316 non-null  object
27  Work_of_casualty                    9118 non-null   object
28  Fitness_of_casualty                 9681 non-null   object
29  Pedestrian_movement                12316 non-null  object
30  Cause_of_accident                  12316 non-null  object
31  Accident_severity                  12316 non-null  object
dtypes: int64(2), object(30)
memory usage: 3.0+ MB

```

The object types are qualitative data while the int type data are quantitative.

Encoding:

It is necessary to replace categorical data with numerical data so that the computer can understand and recognize a pattern to make a prediction.

```
1 #Encoding features
2
3 from sklearn.preprocessing import LabelEncoder
4
5 # List of columns to label encode
6 columns_to_encode = ['Day_of_week',
7 'Age_band_of_driver',
8 'Educational_level',
9 'Driving_experience',
10 'Type_of_vehicle',
11 'Area_accident_occured',
12 'Road_surface_type',
13 'Road_surface_conditions',
14 'Light_conditions',
15 'Weather_conditions',
16 'Type_of_collision',
17 'Vehicle_movement',
18 'Pedestrian_movement',
19 'Cause_of_accident',
20 'Accident_severity']
21
22 # Initialize LabelEncoder
23 label_encoder = LabelEncoder()
24
25 # Apply label encoding to each column
26 for column in columns_to_encode:
27     road2[column] = label_encoder.fit_transform(road2[column])
```

FEATURE SCALING

One of the problems that the algorithm often faces is that it gets biased towards larger values. To avoid this problem, we scale the numerical data to bring it in a small range.

```

2
3 x = road2.drop(['Accident_severity'], axis = 1)
4 y = road2[['Accident_severity']]

9] 1 from sklearn.preprocessing import StandardScaler
    2 scaler = StandardScaler()
    3
    4 scaler.fit(x)
    5 x = scaler.transform(x)

10] 1 '''from sklearn.model_selection import train_test_split
     2 xtrain, xtest, ytrain, ytest = train_test_split(x,y,test_size=0.30, random_state=0, stratify = y)'''
     3
     4 'from sklearn.model_selection import train_test_split\nxtrain, xtest, ytrain, ytest = train_test_split(x,y,test_size=0.30, random_state=0, stratify = y)'

11] 1 from sklearn.model_selection import train_test_split
     2 xtrain, xtest, ytrain, ytest = train_test_split(x,y,test_size=0.30)

12] 1 from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score, ConfusionMatrixDisplay, classi

```

DATASET SPLITTING

The data from the dataset was split into 70% for the training model, and 30% for the testing model.

MODEL TRAINING

In total, six machine learning models have been used for training in both scenarios of before oversampling and after oversampling. They are- Naive Bayes, Decision Tree, K-nearest Neighbors, Random Forest, Support Vector Machine, and MLP.

MODEL TESTING

After training, the six mentioned models were tested.

Before oversampling, the results of performance metrics like accuracy, precision, recall, F1-Score are given below-

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)
-------	--------------	---------------	------------	--------------

Naive Bayes	82.74	76.60	82.74	76.59
Decision Tree	74.09	74.97	74.09	74.52
KNN	80.49	74.93	80.49	76.98
Random Forest	84.27	80.66	84.27	78.30
SVM	84.05	85.50	84.05	76.80
MLP	83.57	76.46	83.57	79.00

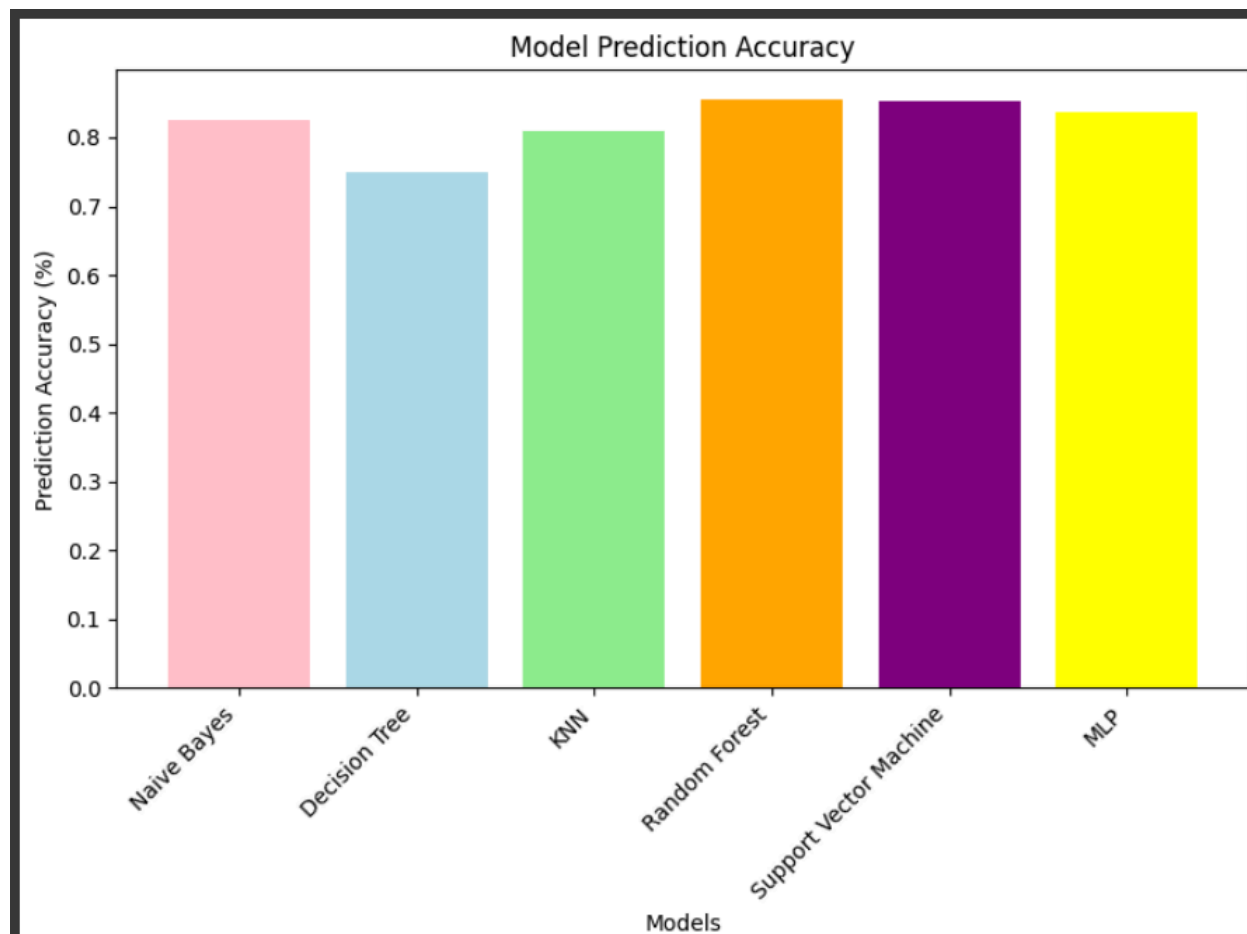
After oversampling, the results of performance metrics like accuracy, precision, recall, F1-Score are given below-

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)
Naive Bayes	43.93	43.70	43.93	42.15
Decision Tree	92.91	93.69	92.91	92.80
KNN	87.80	89.31	87.80	87.46
Random Forest	97.69	97.74	97.69	97.69
SVM	69.71	68.96	69.71	69.12
MLP	78.54	78.29	78.54	78.37


MODEL SELECTION/COMPARISON ANALYSIS

Bar chart showcasing prediction accuracy of all models-

Before oversampling-



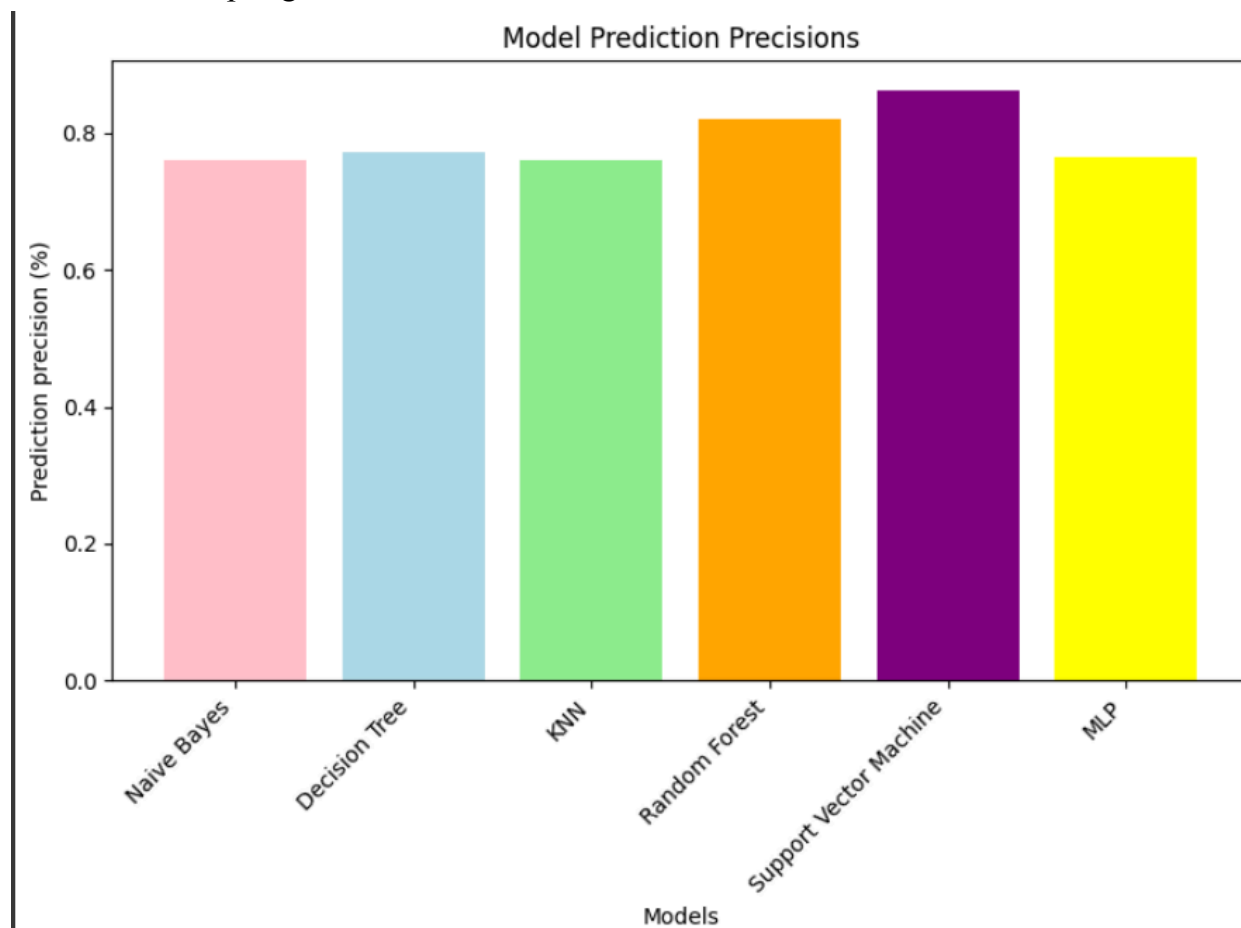
After oversampling-




It can be seen that the accuracies before oversampling were above 0.70 for all the models. After oversampling, the accuracy of Naive Bayes, Support Vector Machine, MLP dropped to below 0.50, 0.70, 0.80 respectively, but on the other hand the accuracies of Decision Tree, Random Forest were boosted to above 0.90, and KNN jumped to 0.87.

Precision comparison of each model:

Before oversampling-




After oversampling-



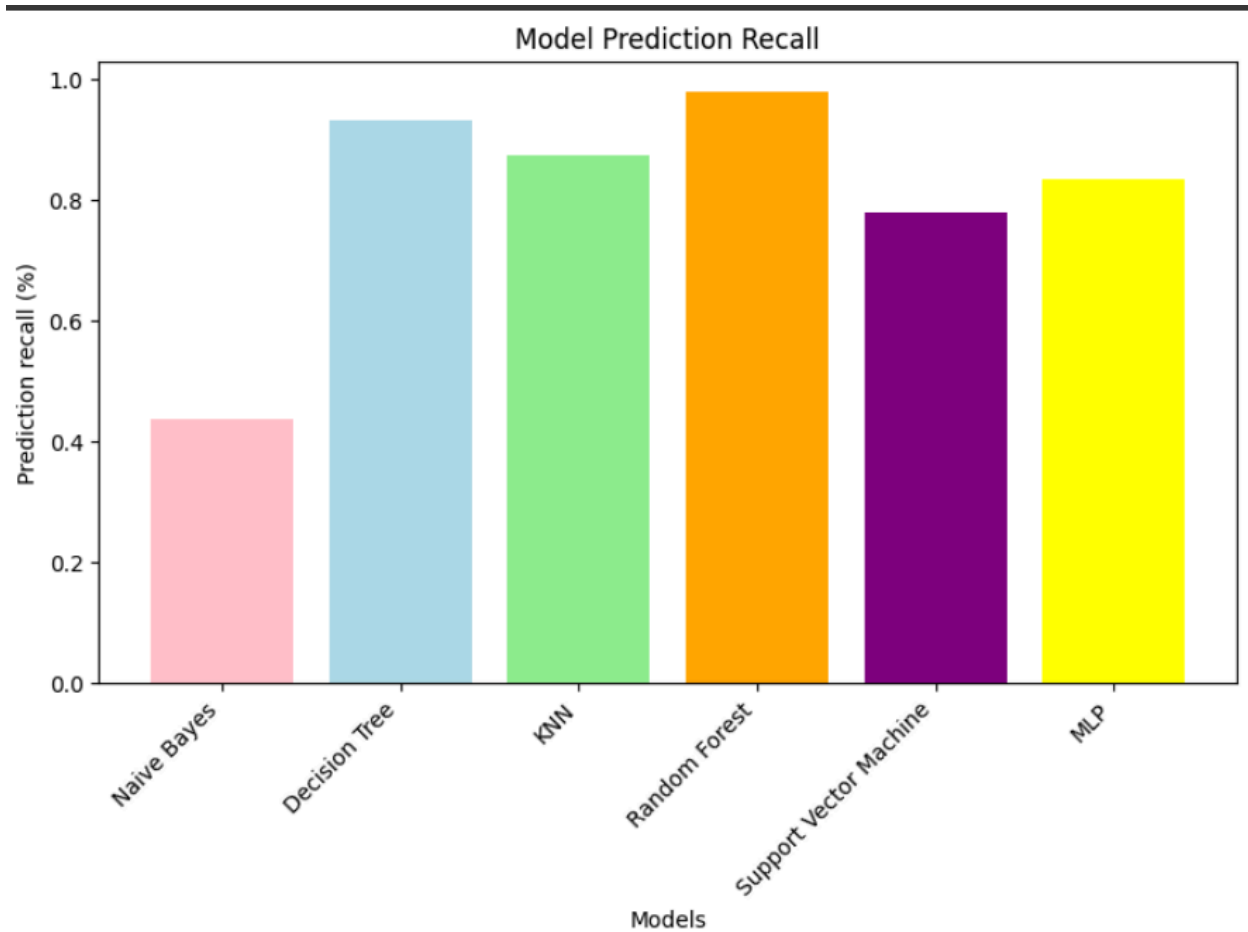
It is evident that before oversampling, the recall for each of the six models is above 0.70. But after oversampling, the precision of Naive Bayes and Support Vector Machine decreased significantly to below 0.50 and 0.70, respectively. While the other four models were able to reach precision levels of around 0.78 (MLP), 0.87 (KNN) and above 0.90 (Decision Tree, Random Forest).

Recall comparison of each model:

Before oversampling:



After oversampling:




It is evident that before oversampling, the recall for each of the six models is above 0.70. But after oversampling, the recall of Naive Bayes, Support Vector Machine, MLP decreased significantly to below 0.50, 0.70, and 0.80 respectively. While the other three models were able to reach recall levels of around 0.87 (KNN) and above 0.90 (Decision Tree, Random Forest).


Confusion Matrix:

Before oversampling:


Naive Bayes:





Decision Tree:




KNN:




Random Forest:**SVM:**




MLP:




After oversampling:


Naive Bayes:**Decision Tree:**




KNN:




Random Forest:



SVM:



MLP:



CONCLUSION

A total of six machine learning algorithms were run, results before oversampling and after oversampling were recorded. Obtained results were accuracies, precisions, recall and F1-score of above 0.70 from scenarios of before oversampling. After the sampling the results of accuracy recorded varied significantly as three of the models were now able to reach 0.90 excluding Naive Bayes, SVM, MLP which performed worse than the other models.

As a result of the insights gained from this study, machine learning in road accident severity assessment can be implemented practically in the future. This development holds substantial ramifications in terms of enhancing global public awareness. In essence, AI holds promise in revolutionizing road safety measures and ensuring the fundamental human right to safe mobility. Through the proactive use of AI technologies, we can strive towards a future where road accidents are minimized, and every individual has the opportunity to travel safely.